10.3.4 Expiring annotations

Annotations for all application groups are kept in a single application group data table, which allows the expiration of annotations to be controlled at a system-wide level. The Life Of Annotations field setup is on the System Parameters General tab. Annotations can be set to never expire or to expire after *N* days. After the number of days (*N*) passes and **ARSMAINT** is run, Content Manager OnDemand removes the annotation.

10.4 Reloading data

If you are migrating data by unloading and then reloading the data, you need to determine your future expiration policy.

Reloading to change the expiration type

For example, if your current expiration policy is set to `Storage Manager` but you later want to perform holds on the data, during the migration process (when you create the application group and before you load any data), change your expiration policy from Storage Manager to `Load`.

When you use the Enhanced Retention Management feature with Content Manager OnDemand or IBM Enterprise Records (formerly IBM FileNet Records Manager), Content Manager OnDemand must be in complete control of expiration processing. Therefore, if you are using Tivoli Storage Manager or OAM, you must disable the ability for either of these storage managers to expire data.

Also, you can use Enhanced Retention Management and Content Federation Services for Content Manager OnDemand only with application groups with an expiration type of `Load`. For those application groups with expiration types of `Document`, `Segment`, or `Storage Manager`, utilities exist to convert these application groups to an expiration type of `Load`.

Consider engaging IBM Lab Services to provide these services.

Reloading ad hoc stored documents

If you choose not to take advantage of the ability of Content Manager OnDemand to aggregate documents but instead you choose to load documents ad hoc by using the storeDocument Java API, StoreDoc Object Linking and Embedding (OLE) API, or CommonStore, you must migrate the data later.

If you choose not to take advantage of the ability of Content Manager OnDemand to aggregate documents into 10 MB storage objects, this decision might result in millions of small objects that are stored in your storage manager, which might cause the storage manager to experience performance problems when it migrates these small objects to tape.

Note: Consider aggregating these smaller objects into larger objects for performance reasons.

For you to aggregate all of these tiny objects into larger objects after they are stored individually requires that you retrieve and reload them as larger objects. You might want to engage IBM Lab Services to assist you with this task.

Another option is to not migrate objects to tape, but to use another random access hardware device instead.